Exploring LLM Capabilities in Abstract Art Analysis

Anthony Battaglia Thomas Lohman Dan Schrage
University of Notre Dame
Department of Computer Science

December 12th, 2025

1 Introduction

In art analysis, where meaning is subjective and ambiguous, the central challenge is aligning a model's reasoning with human interpretation. Successfully developing models that can accurately reflect human emotions and reasoning in response to abstract art would represent a meaningful step toward aligning AI systems with nuanced and inherently subjective aspects of human cognition. Such alignment is essential for building more human-centered models capable of operating in creative, emotional, and interpretive domains (mental health, art critics, and political discourse, etc.).

However, abstract art poses unique challenges for computational modeling. Its interpretation is fundamentally subjective [11] and often intentionally ambiguous [2], making consistent annotation and evaluation difficult. In addition to these conceptual challenges, progress in this area is limited by dataset availability. The ArtEmis dataset [1], a structured collection of abstract artworks paired with human-annotated emotional labels and natural-language explanations, gives a good starting point for studying human–model alignment in this domain.

Our initial objective was to develop a framework for aligning large language models (LLMs) with human reasoning about abstract concepts in art using the ArtEmis dataset. However, we identified a gap in the existing literature: prior work has not examined whether contemporary multimodal LLMs already possess a latent ability to align with human emotional reasoning in art analysis (without task-specific fine-tuning). This observation motivated a more fundamental research question—*to what extent do current open-weight multimodal LLMs exhibit human-aligned reasoning about abstract art under in-context learning alone?*

To investigate this question, we evaluate several open-weight multimodal LLMs available through HuggingFace for human alignment on abstract art analysis. First, we assess how closely the models' predicted emotional classifications align with human annotations. Second, we evaluate the degree to which the models' natural-language explanations resemble human reasoning for the associated emotional responses. Each model is prompted using the same instructions provided to human annotators in the ArtEmis data collection process and is evaluated under varying in-context learning (ICL) regimes, including 0-, 1-, 2-, and 4-shot settings.

Through this study, we aim to establish an empirical baseline for the human alignment capabilities of current multimodal LLMs in abstract art analysis. Our results are intended to inform whether fine-tuning is necessary for achieving human-aligned abstract

reasoning in this domain and to surface systematic differences between human and model interpretations of abstract art. Identifying such patterns may provide broader insights into the strengths and limitations of LLM-based reasoning in subjective and interpretive tasks.

2 Related Work

In this section, we review prior work at the intersection of computer vision, natural language generation, and affective computing. We first discuss traditional approaches to emotion recognition from images, then turn to affective captioning datasets and models, and finally highlight recent advances in large-scale vision-language models and large language models (LLMs).

2.1 Emotion Recognition from Visual Content

Early research on computational aesthetics and emotion recognition focused on predicting discrete affective states directly from visual features. For example, a system might classify an image of a smiling person with the caption *"happy"*, relying only on low-level visual cues such as facial expressions, color palettes, or composition. Works such as Machajdik and Hanbury [9] and Yanulevskaya et al. [15] used psychology-inspired or holistic image features to classify the emotional content of images. With the rise of deep learning, larger-scale benchmarks such as the image emotion recognition dataset by You et al. [17] and subsequent deep neural approaches [7] pushed the field toward robust recognition of basic emotions. However, these methods largely framed emotional responses as a classification task, without addressing the linguistic explanations or interpretive processes that shape human engagement with art. As such, we hope to build upon the existing work by more deeply exploring how LLMs engage with art and how they would describe their interpretive process for classifying emotional responses from the art.

2.2 Affective Captioning and the ArtEmis Dataset

Image captioning has traditionally focused on objective descriptions of natural images. Benchmarks such as MS-COCO [4], Flickr30K [18], Conceptual Captions [12], and Visual Genome [8] emphasize factual descriptions of objects, scenes, and relationships (e.g., "a man walking on the trail"), rather than subjective or emotional interpretations. Early neural captioning systems such as Show-and-Tell [13] and Show-Attend-and-Tell [14] demonstrated the ability to generate fluent, grounded captions, later extended by transformer-based models such as the Meshed-Memory Transformer [6]. These approaches focused on aligning visual features with natural language, enabling models to describe objects and scenes with increasing accuracy and contextual awareness.

To move beyond objectivity, SentiCap [10] introduced explicit sentiment supervision for caption generation, later extended by You, Jin, and Luo [16]. However, these works primarily injected binary sentiment (positive/negative) without providing detailed explanatory language.

The ArtEmis dataset [1] marks a significant milestone by linking over 450,000 artwork annotations with corresponding emotion labels and detailed free-form explanations. ArtEmis highlights subjective, metaphorical, and affective dimensions of art interpretation. This work offers a benchmark for generating captions in addition to emotionally grounded explanations.

Our hope in this work is to continue the efforts in affective captioning by doing more fine-grained sentiment analysis than just positive or negative, instead exploring some of the more specific emotions encom-

passed by those two general sentiments. By using the ArtEmis dataset we can evaluate some specified sentiments in order to see how closely aligned LLM generated affective captioning is to human written affective captioning.

2.3 Large Language Models for Interpretation

Large language models such as GPT-3 [3] and PaLM [5] have shown remarkable capacity for contextual reasoning and explanation generation. When coupled with visual encoders, these models can move beyond surface-level captioning to provide interpretive or evaluative language. For instance, work in multimodal reasoning has shown that LLMs can generate analogies, abstract descriptions, or personalized interpretations given appropriate prompts or fine-tuning. These prior findings motivate our project, which situates LLMs within the context of art interpretation and affective explanation.

3 Methodology

This section outlines the methodological framework used to assess how multimodal models align with human emotional interpretations of visual art under varying in-context learning (ICL) conditions. We begin by describing the dataset and its properties followed by defining the experimental tasks and evaluation criteria. We conclude with a description of the prompting strategy and execution pipeline used across models and shot configurations.

3.1 Dataset Description

We use the ArtEmis dataset collected by Achlioptas et al. in their work ArtEmis: Affective Language for Visual Art[1]. The tabular data they collected was obtained by our team on October 7th, 2025 by filling out a form on their website `https://www.artemisdataset.org/`. The dataset contains 80,031 images and 454,684 image explanations/descriptions. The number of descriptions per painting has a long-tailed distribution (with 5 and 6 descriptions per painting being the most common). When filtering for abstract art, the dataset contains 7,991 images and 47,099 image explanations/descriptions (the data for just abstract art has a similar shape as the overall data which can be seen in Figure 1).

The dataset does not explicitly split its art based on whether it is abstract or not, but every data point has a categorical art_style feature. These art styles were then classified into abstract and concrete based on online research about what artists and art galleries classify these categories as. The data has 5 features: the art style, a unique identifier for the painting, an emotion classification, a description of what evoked said emotion, and a count of how many descriptions there are for that image. The possible emotional classifications are: amusement, anger, awe, contentment, disgust, excitement, fear, sadness, and something else (where something else is a catch all category for artwork that leaves the viewer unsure of what they are feeling when they look at it). Interestingly, if we map out the percent of emotional classifications that fall under "something else" and then order the art styles by that percentage, they naturally segment into their abstract and concrete categories[1] as seen in Table 1. This suggests that abstract art presents interpretive challenges for both AI systems and human annotators, indicating a potential role for AI systems in supporting or complementing human evaluation.

[1]Cubism is actually known as a semi-abstract art form and Pop Art, despite being considered concrete art, is art made using pop culture, which is an abstract concept. It is quite interesting that these two define the boundaries and are also quite similar to each other, since they both use abstract and non-abstract components.

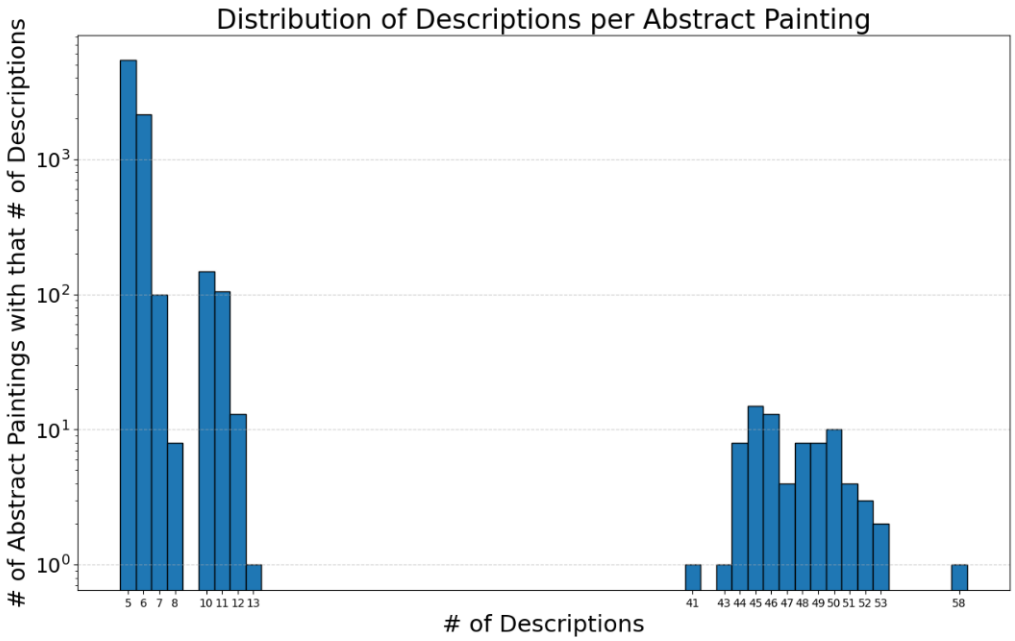

Figure 1: Distribution of Descriptions per Abstract Painting

3.2 Experiment Definition

Given an artwork image and a variable number of in-context examples, our experiments evaluate model performance on two related tasks:

1. **Emotion classification**: predicting the emotion evoked by the image.

2. **Explanation generation**: producing a short natural-language explanation that references visual evidence (e.g., objects, colors, or composition) supporting the predicted emotion.

The first task a multi-class classification problem and is used to assess the alignment between model predictions and the range of human emotional responses captured in the ArtEmis dataset. Because different viewers may experience different emotions when observing the same artwork, there is often no single ground-truth emotion label. Accordingly, a model's prediction is considered aligned if it matches *any* of the emotion labels provided by human annotators for that artwork, rather than requiring agreement with all annotations.

The second task is a text generation problem. For each image, the model first predicts an evoked emotion and then generates a brief explanation describing why the image conveys that emotion (with reference to relevant visual features). As with emotion classification, explanation alignment is evaluated by comparing the model-generated explanation to the set of human-written explanations for the same artwork. Given the diversity of valid human interpretations, we assess whether the model's explanation aligns with at least one human explanation, rather than attempting to enforce agreement with all available descriptions.

Table 1: Percent of Emotion Classifications Selected to be "Something Else"

Art Style	Percent
Minimalism	32.78%
Color_Field_Painting	26.06%
Synthetic_Cubism	25.86%
Analytical_Cubism	22.44%
Action_painting	22.14%
Abstract_Expressionism	20.16%
Cubism	16.39%
Pop_Art	16.14%
Expressionism	12.43%
Early_Renaissance	12.06%
High_Renaissance	12.01%
Art_Nouveau_Modern	11.87%
New_Realism	11.31%
Naive_Art_Primitivism	11.27%
Ukiyo_e	10.98%
Mannerism_Late_Renaissance	10.88%
Realism	10.79%
Symbolism	10.61%
Contemporary_Realism	10.53%
Post_Impressionism	10.45%
Baroque	10.25%
Fauvism	10.12%
Rococo	10.08%
Northern_Renaissance	10.08%
Romanticism	8.88%
Pointillism	8.82%
Impressionism	8.05%

3.3 Experiment Pipeline

Our experimental pipeline consists of dataset partitioning, prompt construction with in-context demonstrations, model inference, structured output parsing, and embedding-based analysis of generated explanations.

We first construct a held-out test set by randomly sampling 100 data points from the abstract art subset of the ArtEmis dataset. These test examples are fixed across all models and in-context learning (ICL) configurations to ensure comparability. The remaining data points (all abstract artworks not included in the test set) form the demonstration pool from which in-context examples are sampled.

For each test example and ICL setting (0-, 1-, 2-, and 4-shot), demonstration examples are randomly selected from the demonstration pool. Each demonstration consists of an image paired with its corresponding human-annotated emotion label and explanation. These demonstrations are inserted into the prompt immediately following the task instructions, maintaining a consistent prompt structure across all experiments. The test image is then appended as the final query instance.

The base prompt used in our experiments was derived from the supplemental materials provided on the ArtEmis project website. This prompt is itself an adaptation of the instructions originally presented to human annotators during the ArtEmis data collection process via a crowdsourcing platform. To ensure compatibility with each evaluated multimodal LLM, minor model-specific adaptations were required. However, the overall structure, task instructions, and response requirements remained consistent across models. A representative version of the prompt template is shown in Appendix A.

Each multimodal LLM is run in inference-only mode using the constructed prompt. Models are instructed to output their response in a structured JSON format of the form {`emotion`: *label*, `utterance`: *explanation*}. The reason we added the JSON formatting restriction into the prompt was twofold: it gave us a structure for showing the models the ICL examples and it made parsing the models' outputs' quick and reliable.

To allow quantitative comparison between model-generated explanations and human-written explana-

tions, we embed the generated explanations using the `BAAI/bge-base-en-v1.5` sentence embedding model. These embeddings are used to compute semantic similarity between model outputs and corresponding human explanations, allowing us to assess alignment at the level of explanatory content rather than surface lexical overlap.

4 Results

In this section, we present the outcomes of our experiments evaluating the ability of several large language and vision-language models to classify emotions in artworks and generate affective descriptions. We report results for both zero-shot and few-shot in-context learning scenarios, comparing model predictions against human annotations. The analysis focuses on two main dimensions: the accuracy of emotion classification and the semantic alignment of model-generated explanations with human-written descriptions. By examining these metrics, we aim to quantify how closely each model mirrors human judgment and interpretation of visual art.

4.1 Evaluation Metrics

To evaluate each model's performance on the emotion classification task, we computed accuracy. Accuracy was chosen to quantify the overall proportion of correctly classified instances. For the sake of this experiment we defined a correctly classified instance as any instance where the model predicted an emotion that matches any of the human annotators emotions for that painting. For example, if a painting has 5 annotations and for those 5 annotations the human responses for emotion were *[sadness, disgust, sadness, something else, anger]* and the model classified the image as *anger*, this would count as an accurate classification because it was aligned to a human response. If the model instead classified the painting as *fear*, this would count as an incorrect classification because it was aligned with none of the human annotators.

To evaluate the alignment between model-generated explanations and human-written explanations, we measure semantic similarity using cosine similarity in a learned sentence embedding space. Both the model-generated explanation (`utterance`) and the corresponding human explanations for each painting are embedded using the `BAAI/bge-base-en-v1.5` embedding model, which is designed to capture semantic similarity between short and medium-length texts.

For each test instance, cosine similarity is computed between the model-generated explanation and each available human explanation associated with the same painting. The highest cosine similarity score is then selected as the explanation alignment score for that instance. This "best-match" strategy is motivated by the fact that ArtEmis provides multiple valid human explanations per artwork, reflecting diverse but equally plausible interpretations. As there is no single canonical reference explanation, selecting the maximum similarity allows us to assess whether the model's explanation aligns semantically with at least one human interpretation rather than penalizing alternative yet reasonable descriptions. Overall explanation alignment is reported as the mean cosine similarity across all test instances for each model and in-context learning setting.

4.2 Baselines

Since our experiment is testing the human alignment of these models, we decided that a good baseline to test against would be human-human agreement. Since each painting had at least 5 classified emotions and affective captions for those emotions, we calculated human-human agreement by iterating through each painting in the test set and for each emotion-caption

pair comparing it against all of the other emotion-caption pairs for that painting. Then, in the same way that we did for the models, we scored the human written emotion-caption pair by seeing if its emotion matched any of the other human selected ones and we scored the caption by having it match the caption that shared the highest cosine similarity score with itself.

4.3 Experimental Results

Table 2 reports emotion classification accuracy and explanation alignment, measured via cosine similarity, across different in-context learning (ICL) settings for each evaluated model. Some consistent trends were observed across models and number of evaluations.

As Table 2 shows, there is no significant change in any of the models' performance in affective captioning. For the InternVL3-8B and Llama-3.2-11B-Vision-Instruct models there is a case that their emotion classification was improved via few-shot prompting, but for the other two models no such claim can be made. All of the models were identifiably less aligned with humans on the emotion classification task than other humans were with themselves, which shows that there is room for further improvement in that category. However, for the human-human agreement on the description generation, all of the models had similar level of agreement with InternVL3-8B being only .001 less aligned, an amount that is experimentally insignificant.

Figures 2 and 3 visualize the distribution of emotion classifications made by the LLMs. The two figures report the same data, however, Figure 2 groups the data by model and Figure 3 groups the data by number of shots.

One of the interesting things to note is how positive all of the models seem to be. The most common emotion classifications across the board were awe and contentment with excitement trailing not too far behind. Part of this is likely due to the fact that most LLMs are trained to please and as such they are predisposed to not be too harsh of critics.

Table 2: Model Performance Across Shot Settings

Shot	Accuracy	Cosine Similarity
LLaVA-OneVision-1.5-8B-Instruct		
0-shot	**0.600**	**0.688**
1-shot	0.560	0.685
2-shot	**0.600**	0.679
4-shot	**0.600**	0.672
InternVL3-8B		
0-shot	0.560	**0.693**
1-shot	0.490	0.692
2-shot	**0.630**	0.685
4-shot	0.620	0.689
Qwen3-VL-8B-Thinking		
0-shot	0.600	0.677
1-shot	**0.640**	0.677
2-shot	0.630	**0.683**
4-shot	0.620	0.679
Llama-3.2-11B-Vision-Instruct		
0-shot	0.590	0.676
1-shot	0.490	0.670
2-shot	0.550	0.669
4-shot	**0.640**	**0.682**
Other Human Annotators		
0-shot	**0.719**	**0.694**

Another very interesting thing to note is that in the 0-shot experiments all of the models classified "awe" at a higher rate than the other emotions. It almost makes us wonder if they are marveling at human creativity and the beauty of creation and expression. Perhaps the reason this "awe" decreases as the shots increase is because the human examples almost demystify the art in a way, giving the models an example of how to think about the art and giving way for the models to process the art on a new level. However, more likely than not is just that in the models' train-

ing data there is something that links art to being awe inspiring. After all, in most history textbooks and sources on the internet ancient statues and paintings are usually talked about in a reverent manner. Regardless of where it comes from, we do still believe though that it is interesting to note how the models tend to default to the classification of awe when given nothing else to guide it.

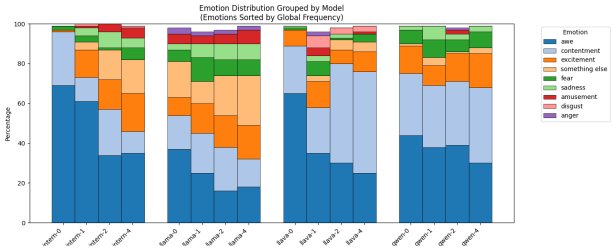

Figure 2: Emotion Distribution Grouped by Model (Emotions Sorted by Global Frequency)

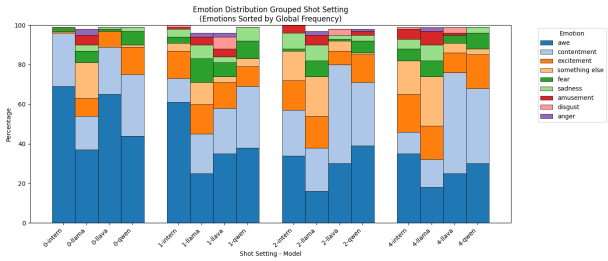

Figure 3: Emotion Distribution Grouped by Shot Setting (Emotions Sorted by Global Frequency)

5 Conclusions

In this work, we investigated whether open-weight multimodal large language models (LLMs) can align with human emotional interpretations of abstract art through in-context learning alone. Our results show that while these models demonstrate meaningful zero-shot performance, increasing the number of in-context demonstrations provides little to no (and inconsistent) improvements in both emotion classification accuracy and explanation alignment. This suggests that in-context learning by itself does not provide sufficient feedback to reliably align model behavior with human reasoning in this highly subjective domain.

At the same time, the relatively strong zero-shot performance observed across models suggests the presence of some latent capacity in pretrained multimodal LLMs to reason about emotional content in abstract art, even without task-specific training. Although this capability is limited and imperfect, it indicates that pretrained representations may already encode aspects of the reasoning required for abstract art interpretation.

We also find that model-generated explanations are no more aligned with human explanations than the level of agreement observed among different human annotators. This highlights both the inherent ambiguity of abstract art interpretation and the difficulty of using explanation alignment as a strict evaluation target.

Finally, we suspect that the quality and consistency of some human-written explanations in the dataset may have constrained achievable alignment, as a significant subset of annotations were weakly grounded in visual evidence (more on this in the limitations section).

6 Limitations

While our study provides insights into the alignment of large language and vision-language models with human affective responses to art, several limitations should be acknowledged. These relate to the quality and completeness of the dataset, the scope of emotional categories considered, and the computational

resources available for model evaluation. Recognizing these constraints is important for properly interpreting the results and understanding the potential impact on model performance and generalizability.

6.1 Dataset Quality

First, because the dataset was publicly sourced through large-scale crowd annotation, the overall quality of responses is variable. Many participants did not fully adhere to the provided prompts, resulting in descriptions or emotional responses that were incomplete, off-topic, or inconsistent with the intended task. Additionally, a portion of participants demonstrated limited familiarity with the artworks or with artistic terminology, which contributed to vague, superficial, or otherwise low-quality annotations. These factors introduce noise into the dataset and may obscure underlying patterns related to affective response or aesthetic interpretation. Consequently, any models trained on ArtEmis should be interpreted with caution, as these inconsistencies may impact both performance and downstream analyses.

6.2 Excluded Emotional Classification

Another limitation of this study was the exclusion of the disgust emotional category. This was purely an oversight on our part. The ArtEmis team did not provide a clearly indicated comprehensive set of emotions that the participants could select from. They did however, provide examples of utterances for what we thought were all 8 emotional response options. It turns out that there was actually a 9th option that no example utterance was shown for. This impacted our study because it led to disgust never being classified until and unless a demonstration example had it as one of the example emotion classifications in the different shot testing. This likely led to slightly reduced performance on accuracy for all of the tests we performed.

6.3 Hardware and Compute Limitations

Our experiments were constrained by hardware availability, which influenced model selection and learning methodology. The work was conducted primarily using Google Colab Pro, which provides limited computational resources compared to dedicated high-performance computing environments. These limitations prevented extensive fine-tuning of large language or multimodal models, necessitating the use of in-context learning approaches instead. While in-context learning allowed us to leverage pre-trained knowledge without full model training, it may be less effective than fine-tuning in capturing dataset-specific patterns, potentially limiting overall performance.

Hardware constraints also restricted our choice of model size. Due to limited memory and processing capacity, we were unable to employ larger models that might achieve higher accuracy or better generalization, and instead focused on smaller architectures. This choice, while practical, likely influenced both the absolute performance of our models and their ability to fully exploit the ArtEmis dataset. Furthermore, these constraints may affect the reproducibility of results; experiments conducted on higher-end hardware with larger models could yield different outcomes. Consequently, the findings reported here should be interpreted with consideration of these computational limitations.

Lastly, limitations in available computational resources constrained the scale of our evaluation, leading us to restrict the test set to 100 abstract paintings.

7 Future Work

While the current study provides insights into model performance on emotion classification and affective

description generation for artworks, several avenues remain for further exploration. Future work can expand the range of models evaluated, improve the quality and reliability of the dataset, and explore advanced training techniques such as fine-tuning and reinforcement learning with human feedback. By pursuing these directions, subsequent research can deepen our understanding of how models interpret art, enhance alignment with human and expert judgments, and improve the overall robustness and interpretability of model outputs.

7.1 Expand Model Choice

Future work should explore a broader range of models to better understand the effects of architecture and model size on performance with the ArtEmis dataset. In the current study, model selection was constrained by hardware limitations, which restricted experiments to smaller architectures and a limited set of models. Expanding the set of models tested—including larger language models, as well as models with different underlying architectures—could reveal emergent capabilities or patterns not observable with smaller models. Additionally, comparing models across different scales may provide insights into how model size impacts the ability to capture nuanced emotional and aesthetic responses in art. Such experiments would allow a more comprehensive assessment of model generalization, robustness, and potential for downstream applications, while informing best practices for future research in computational art understanding.

7.2 Clean ArtEmis Dataset

The quality of the ArtEmis dataset poses a significant limitation for computational experiments. Many responses in the dataset are incomplete, off-topic, or otherwise fail to follow the provided prompts, while some annotators appear to have a limited understanding of the artworks they evaluated. These issues introduce noise and reduces the reliability of demonstration examples for ICL experiments and training examples for fine-tuning.

A promising direction for future work is to clean and curate the dataset using large language models (LLMs). By automatically identifying and removing responses that do not adhere to the prompt or that are of low quality, it would be possible to create a higher-fidelity version of ArtEmis. Such cleaning could involve scoring responses for relevance, coherence, and completeness, or flagging entries that demonstrate misunderstanding of the artwork. A curated dataset would reduce noise, improve model training, and enable more accurate assessments of model capabilities in interpreting art and generating affective responses.

7.3 Reinforcement Learning with Human Feedback

A promising direction for future work is to combine reinforcement learning with human feedback (RLHF) and expert evaluation to improve model alignment and output quality. In this approach, models would be fine-tuned based on evaluative feedback from experts, such as trained art historians, curators, or art educators, rather than relying solely on crowd-sourced annotations. Expert guidance can provide nuanced judgments regarding style, technique, and emotional expression, allowing models to learn to interpret art in ways that reflect professional standards.

By incorporating experts in the feedback loop, RLHF can help the model refine its responses to better capture subtle aesthetic and affective features, producing outputs that align more closely with expert-level interpretation. This expert-guided training approach would shift model behavior from "general population alignment" toward "expert alignment," increasing the reliability, precision, and sophistication

of model-generated responses. Such a framework not only has the potential to improve performance on curated datasets like ArtEmis but also provides a pathway for evaluating and enhancing model reasoning in complex, subjective domains such as art interpretation.

7.4 Fine Tuning

Another avenue for future work is to perform full fine-tuning of models on a cleaned dataset. While the current study was constrained to in-context learning due to hardware limitations, fine-tuning allows models to adjust their internal weights specifically to the task of art interpretation and emotional response generation. By training directly on high-quality, curated data—potentially filtered through expert review or automated cleaning processes—fine-tuned models may better capture subtle patterns, stylistic nuances, and affective cues that in-context learning alone may miss.

Fine-tuning also enables more control over model behavior, potentially reducing inconsistencies and improving alignment with expert or desired outputs. Combined with approaches such as expert-guided reinforcement learning, fine-tuning can help create models that are not only more accurate but also more interpretable and reliable in complex, subjective tasks like analyzing art and generating emotionally informed responses.

7.5 Internal Alignment

As was noted in the ArtEmis Project[1] models are known to say that they chose or did something for one reason, when in reality they did it for a completely different one. As such, a nice interpretability and transparency measure would be performing future work to see what parts of the image/prompt the model payed attention to when making the emotion classification. These attention weights could then be transformed into heat maps so that a human reviewer could see if the parts of the painting the model payed attention to when making the classification are the same parts of the painting that the model references in its textual explanation of that classification. If it turns out that models are not internally aligned in what they look at to make the classificaiton and how they justify their explanation, we think it would be very interesting to see how a model that is internally aligned performs and what new insights that might generate.

A Prompt Template

```
STEP 1: How does this painting make you primarily feel? Select from the
↪  following: Amusement, Awe, Contentment, Excitement, Fear, Sadness, Anger,
↪  Something-Else

then...

STEP 2: Give a detailed description (at least 7 words) about WHY you feel like
↪  this, based on SPECIFIC details of the painting.

Examples of GOOD descriptions:
  - "the sky looks gloomy and the shadows are scary"
  - "the red marks on the table look like drops of blood" (analogies are
   ↪  good!)
  - "the blue color of the lake contrasts with the orange hats of the men"

Examples of BAD descriptions:
  - Vague descriptions that do not explain WHY you felt like this:
    - "I like its colors" (be more specific)
    - "It is amazing, nice work! great painting" (why is it amazing? be
     ↪  specific)
    - "I don't know what this is" (try and explain what it looks like)

Note: if you selected "Something-Else" also explain how you felt in addition
↪  to why you felt Something-Else.


Here is example {n}:
{demonstration_image_n}
{
    "emotion": "{demo_emotion_n}",
    "utterance": "{demo_utterance_n}"
}


Finally, output your response strictly in the following JSON format:

{
    "emotion": string,
    "utterance": string
}

This is the image that you are to generate a response for:
{test_image}
```

References

[1] Panos Achlioptas et al. "ArtEmis: Affective Language for Visual Art". In: *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*. 2021.

[2] Xavier Alameda-Pineda et al. "Recognizing Emotions from Abstract Paintings using Non-Linear Matrix Completion". In: *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*. 2016, pp. 2624–2633. DOI: 10.1109/CVPR.2016.287. URL: https://openaccess.thecvf.com/content_cvpr_2016/papers/Alameda-Pineda_Recognizing_Emotions_From_CVPR_2016_paper.pdf.

[3] Tom B Brown et al. "Language models are few-shot learners". In: *Advances in Neural Information Processing Systems (NeurIPS)*. 2020.

[4] Xinlei Chen et al. "Microsoft COCO captions: Data collection and evaluation server". In: *arXiv preprint arXiv:1504.00325* (2015).

[5] Aakanksha Chowdhery et al. "PaLM: Scaling language modeling with pathways". In: *Proceedings of the International Conference on Machine Learning (ICML)*. 2022.

[6] Marcella Cornia et al. "Meshed-memory transformer for image captioning". In: *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*. 2020.

[7] Heechul Kim et al. "Building emotional machines: Recognizing image emotions through deep neural networks". In: *IEEE Transactions on Multimedia* (2018).

[8] Ranjay Krishna et al. "Visual Genome: Connecting language and vision using crowdsourced dense image annotations". In: *International Journal of Computer Vision (IJCV)*. 2017.

[9] Jana Machajdik and Allan Hanbury. "Affective image classification using features inspired by psychology and art theory". In: *ACM International Conference on Multimedia*. 2010.

[10] Alexander Patrick Mathews, Lexing Xie, and Xuming He. "SentiCap: Generating image descriptions with sentiments". In: *AAAI Conference on Artificial Intelligence*. 2016.

[11] David Melcher and Francesca Bacci. "Chapter 10 - Perception of emotion in abstract artworks: A multidisciplinary approach". In: *The Fine Arts, Neurology, and Neuroscience*. Ed. by Stanley Finger et al. Vol. 204. Progress in Brain Research. Elsevier, 2013, pp. 191–216. DOI: https://doi.org/10.1016/B978-0-444-63287-6.00010-5. URL: https://www.sciencedirect.com/science/article/pii/B9780444632876000105.

[12] Piyush Sharma et al. "Conceptual captions: A cleaned, hypernymed, image alt-text dataset for automatic image captioning". In: *Proceedings of the 56th Annual Meeting of the Association for Computational Linguistics (ACL)*. 2018.

[13] Oriol Vinyals et al. "Show and tell: A neural image caption generator". In: *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*. 2015.

[14] Kelvin Xu et al. "Show, attend and tell: Neural image caption generation with visual attention". In: *International Conference on Machine Learning (ICML)*. 2015.

[15] Victoria Yanulevskaya et al. “Emotional valence categorization using holistic image features”. In: *IEEE International Conference on Image Processing (ICIP)*. 2008.

[16] Quanzeng You, Hailin Jin, and Jiebo Luo. “Image captioning at will: A versatile scheme for effectively injecting sentiments into image descriptions”. In: *arXiv preprint arXiv:1801.10121* (2018).

[17] Quanzeng You et al. “Building a large scale dataset for image emotion recognition: The fine print and the benchmark”. In: *ACM International Conference on Multimedia*. 2016.

[18] Peter Young et al. “From image descriptions to visual denotations: New similarity metrics for semantic inference over event descriptions”. In: *Transactions of the Association for Computational Linguistics*. 2014.